Is the latency saving reproducible across repeated measurements?

32 shared dev questions × 3 methods × 3 observations; trial 0 is historical, trials 1–2 are new

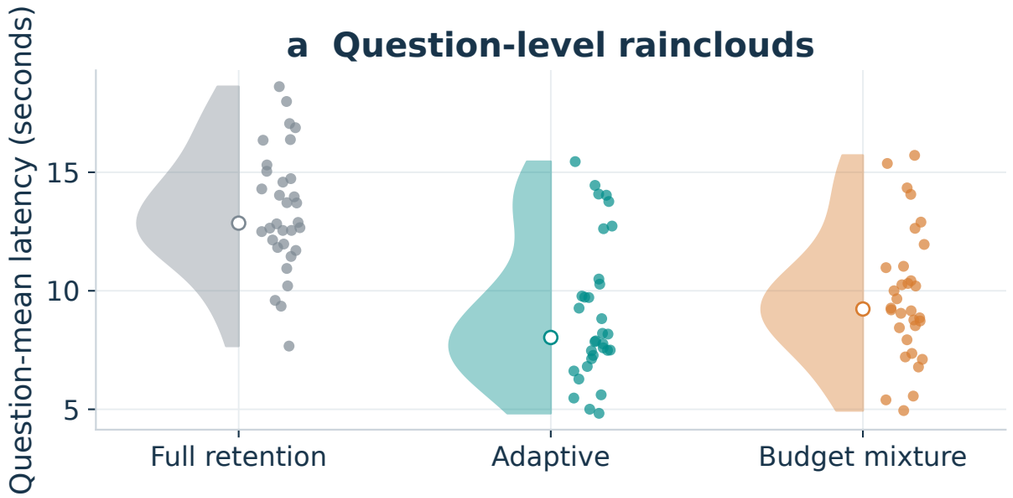

b Mean seconds by trial

	Historical 0	New 1	New 2
Full retention	13.35	13.39	13.39
Adaptive	9.06	9.07	9.08
Budget mixture	9.74	9.76	9.76

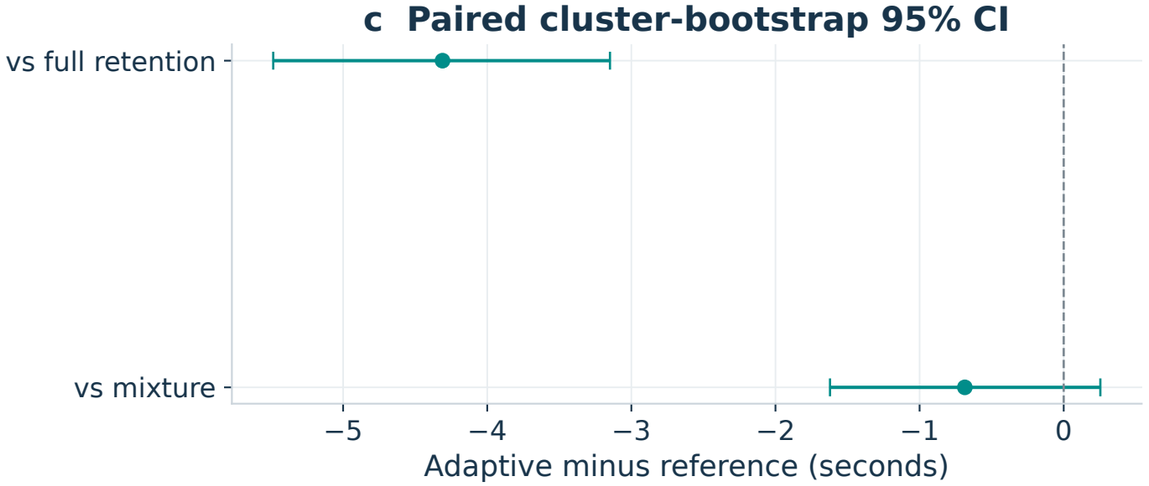

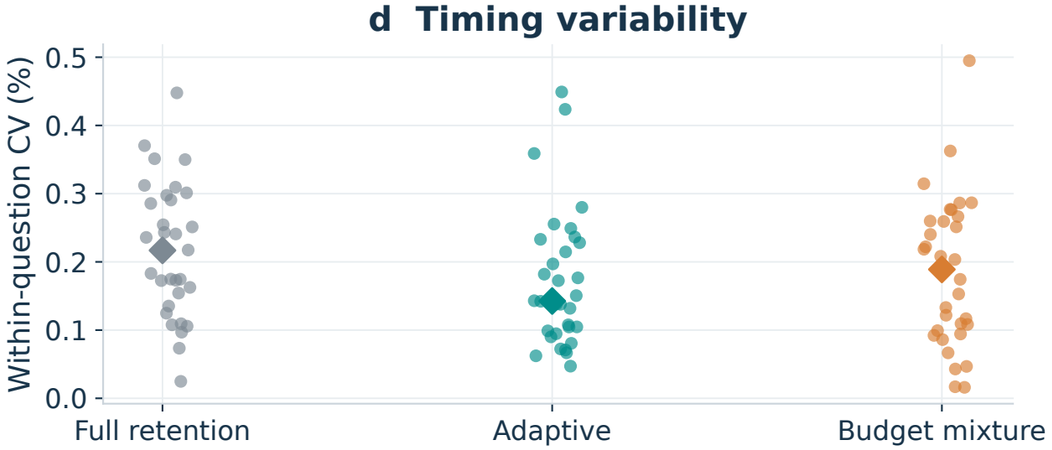

Cluster unit = question, n=32; repeats are not independent questions/training seeds. Token-sequence inconsistent query-methods: 0/96.